

Mesh partitioning and node ordering for FESOM2/Kokkos: measurements on CPU and GPU, and a procedure for adopting a new partition

Campaign report

12 August 2026 — Levante (128-core CPU nodes, 4×A100 GPU nodes; GH200 testbed)

TL;DR

- Repartitioning makes the GPU model step **8–20 % faster** (DARS on 64 GPUs: -14.3% with the disturbance confined to sea-surface height, -19.7% with the largest disturbance measured in the campaign). The gain comes from one METIS option, MINCONN, which FESOM’s `PartGraphRecursive` call silently ignores; switching to `PartGraphKway` is a few lines.
- Any repartitioning perturbs the solution; we report each candidate’s perturbation against the spread of ordinary seed re-rolls of the stock recipe and interpret it (Section 6). Most winners sit inside or below that spread. At three of the four GPU points the winner exceeds it in exactly one statistic, the SSH rms — at DARS traceably an artefact of the barotropic solver’s stopping tolerance, not a physics bias.
- CPU runs gain **4–8 %**, from external engines (KaMinPar, Mt-KaHyPar) and load-balance slack instead; the two backends need different partitions of the same mesh.
- **Some new partitions crash the model** — eight caught, one produced by the unmodified stock recipe. Nothing computable from the partition predicts which. Every new partition must complete 3,000 steps at its target rank count before use; on failure, regenerate with a new random seed.
- Hilbert renumbering pays only on CORE2 (-5.8% combined with repartitioning).

1. Summary

FESOM2 reads its domain decomposition from files, so a better partition can be adopted without touching model code or affecting any other configuration. We generated several hundred candidate partitions of four meshes with METIS under settings FESOM does not currently use, with three external partitioning engines, and with two node renumberings, and timed them against the shipped partitions on both backends of the C++/Kokkos port.

Three results carry the report. First, the shipped partitions cost the GPU backend up to 20 % of the model step, and most of that is recovered by a single METIS option, MINCONN, which no FESOM partition has ever had active: the partitioner calls `METIS_PartGraphRecursive`, which accepts and ignores it. Each gain is reported together with the disturbance the new partition causes, measured against the spread of ordinary seed re-rolls of the shipped recipe: the two largest recommended gains (DARS -14.3% , NG5 -9.8%) exceed that spread in the SSH rms alone, and the largest measured gain (-19.7% on DARS) carries a temperature-field disturbance that should be checked with a long twin run before production (Section 6). Second, the CPU backend gains less (4–8 %) and gains it from different settings, namely external engines and load-balance slack; the two backends are best served by different partitions of the same mesh. Third, and unexpectedly, a small fraction of newly generated partitions cause the model to diverge, regardless of which settings produced them. Eight such partitions were caught during the campaign, one of them generated by the unmodified shipped recipe. Because no property computable from the partition itself identifies the dangerous ones, any partition intended for production must first complete a stability run at its target rank count; regenerating a failed partition with a different random seed has so far always produced a sound replacement.

Meshes: CORE2 (127k surface nodes), fArc (638k), DARS (3.16M), NG5 (7.40M). All experiments are cold starts under JRA55 (1958) forcing from PHC climatology, at the largest time step at which

a cold start of the given mesh is known to be stable: 1800s on CORE2, 900s on fArc, 120s on DARS, 180s on NG5.

2. Candidates

METIS settings. Switching the METIS call from `PartGraphRecursive` to `PartGraphKway` makes three previously ignored options effective: the communication-volume objective, contiguity of the parts (`CONTIG`), and `MINCONN`, which minimises the largest number of neighbouring subdomains any part has. We swept these together with load-balance slack (`UFACTOR` of 10, 30 and 100), a family of vertical weights $w = a + nlev$ with a between 0 and 100, two-level hierarchical partitions, and the METIS random seed. The Recursive call is not an oversight in the original code: `fort_part.c` documents that it produced the better partition on a test mesh, and notes which options it ignores. The comparison, however, predates the GPU backend.

External engines. KaHIP, KaMinPar and Mt-KaHyPar, each given the same node-weighted dual graph and the same weight sweep.

Node renumbering. Reordering the mesh files themselves along a Hilbert space-filling curve, and alternatively by reverse Cuthill–McKee, to improve memory locality. Unlike a partition, a renumbering changes the mesh files, so every reference solution pinned to the old numbering must be re-established.

A second architecture. The best GPU settings were re-timed on DKRZ’s GH200 nodes (dolpung), whose interconnect cannot register GPU memory; halos there are staged through pinned host buffers, which prices communication differently than the A100 fabric does.

3. Measurement procedure

Timing runs place all candidates in a single job allocation and interleave repetitions, so that node-to-node and time-of-day variation affects every candidate alike. Each run integrates 300 steps; we quote the best of N repetitions, as a percentage of the shipped partition’s time per step. The mesh-definition files are verified byte-identical across candidates before the job starts, so the candidates differ in the decomposition only.

A timing result is accepted only after two further tests.

Stability. The candidate must complete 3,000 steps (roughly two model months) at the rank count at which it will be used. This test exists because it fails in practice: a DARS partition that had won its timing comparison by 4.5 % diverged between step 150 and step 3,000 while the shipped partition ran cleanly.

Accuracy. Any change of decomposition reorders floating-point reductions, and the resulting round-off differences grow chaotically at fronts, so a new partition cannot be required to reproduce the reference solution. What can be measured is whether it is *distinguishable* from an ordinary repartitioning. We regenerate the shipped recipe with three or more different random seeds (verifying that each is a genuinely different partition), run twenty steps with each, and report the candidate’s temperature, salinity and sea-surface-height rms differences from the reference against the range spanned by those control partitions. That spread is the scale bar for every disturbance quoted in Section 6; quantiles of the difference fields are recorded as context. These are twenty-step cold-start numbers: they measure the detectability of the partition change, not its climate impact (Section 6 discusses the distinction).

Recommended partitions are packaged by a tool that refuses any decomposition lacking a recorded clean 3,000-step run at its exact rank count.

4. Speed

Figure 1 shows the best candidate at each mesh and rank count; Table 1 records the settings and the state of the evidence behind each bar. Every gain that was re-measured over 3,000 steps came back equal to or larger than its 300-step value (Fig. 2b; for example -18.6 to -19.7 % on DARS




Figure 1: Best candidate partitions per mesh and rank count, as a percentage faster than the shipped partition. Gray text under each label is the setting that produced the candidate. All bars passed the 3,000-step stability screen; the fill codes the twenty-step disturbance against the seed-control spread (Section 6): solid = inside or below the spread, single-hatched = above it in the SSH rms only, cross-hatched = above it in temperature or salinity (the grade that warrants a long twin run before production). Mesh colors follow the project’s figure conventions and recur in Figs. 2 and 3.

GPU), so the short timing runs are, if anything, conservative.

mesh	backend, ranks	gain	setting	stability	disturbance vs seed spread
DARS	GPU, 64	−19.7 %	MINCONN+CONTIG+u30	passed	T rms +13 %, T max $\times 1.8-2.5$
DARS	GPU, 64	−14.3 %	MINCONN	passed	ssh rms +11 % only (solver stopping)
NG5	GPU, 64	−9.8 %	MINCONN	passed	ssh rms +8 % only (unexplained)
CORE2	GPU, 4	−8.1 %	MINCONN	passed	below every control
fArc	GPU, 16	−3.6 %	MINCONN+CONTIG	passed	T rms +33 %, ssh +68 %
fArc	CPU, 2048	−7.5 %	Mt-KaHyPar $w=100+nlev$	passed	inside the spread (4 controls)
CORE2	CPU, 512	−5.8 %	Hilbert renumbering + engine	passed	inside the spread
DARS	CPU, 2048	−4.2 %	KaMinPar $w=100+nlev$	passed	below every control (T, S)
CORE2	CPU, 864	−4.9 %	KaMinPar	passed	salinity rms +24 %
CORE2	CPU, 512	−3.8 %	UFACTOR=30 alone	passed	inside the spread
NG5	CPU, 2048	null	—	—	—

Table 1: Best candidates per mesh and rank count, relative to the shipped partition, at the time step given in Section 1. The last column is the twenty-step rms difference from the reference where it exceeds the spread of seed-control repartitionings, with the responsible field; Section 6 interprets each entry.

4.1. What the two backends pay for

Across the nine groups of candidates that were timed, the number of communication partners per rank is the only precomputable quantity that correlates positively with GPU step time. Every volume-like quantity — edge cut, halo bytes, replicated elements — correlates negatively: on the GPU, candidates that ship more data are faster when they talk to fewer neighbours. MINCONN is the METIS objective that acts on partner count, which is why it wins on the GPU and why the win remained invisible while the Recursive call ignored the option.

The size of the GPU gain across meshes follows the fractional reduction in the *maximum* partner count (Fig. 2a); the mean partner count is reduced by a similar fraction at all three points and therefore explains none of the ordering. That the maximum is the relevant statistic is expected, since each step waits for the slowest rank. The relation holds within one machine; Section 4.3 shows that its coefficient does not transfer across interconnects.

On the CPU, no precomputed quantity orders the candidates: for load imbalance, edge cut and partner count alike, the rank correlation with step time straddles zero across the timed groups. CPU candidates therefore have to be timed. What does generalise is that an external engine won at every CPU point with 512 or more ranks, and that relative to MINCONN alone, adding CONTIG and slack




Figure 2: (a) GPU gain against the fractional reduction of the maximum per-rank communication-partner count, annotated with the counts themselves. (b) Gain measured in the 300-step timing run against the same candidate’s gain re-measured over 3,000 steps, for the eight candidates measured both ways; the line marks equality, circles are GPU points, squares CPU, and colors identify the mesh as in Fig. 1.

gains five percentage points on DARS GPU and loses one to two on DARS CPU: the same options move the two backends in opposite directions.

4.2. Node renumbering pays only on CORE2

Hilbert renumbering is worth -1.2 to -2.4 % on CPU and -5.0 % on a single GPU node on CORE2, and combines almost additively with repartitioning (-5.8 % at 512 ranks). The other meshes gain nothing, because their element-gather locality as shipped is already 88 %, against 27.6 % for CORE2. Since a renumbering forces all pinned references to be re-established (Section 2), we recommend it only where that cost has been accepted; the repartitioning gain requires no such decision.

4.3. The second architecture

On a single GH200 node the CORE2 result reproduces slightly above its A100 value (-8.9 against -8.1 % with MINCONN, two independent jobs agreeing within 0.6 percentage points), so the effect is not specific to A100 hardware. It did not, however, survive the move to several nodes (Fig. 3): a prediction, registered before the job ran, that DARS on 64 GH200 ranks would exceed its A100 gain of -18.6 % failed, with the measured effect between -1 and -3 %, and fArc on 4 nodes shows no effect at all at five repetitions. On this machine the per-message cost that MINCONN removes evidently resides in the staged intra-node path, not in the inter-node fabric. Multi-node gains should therefore be measured on each interconnect rather than transferred. The same experiments showed that quoting the best of N repetitions is biased whenever the reference partition is noisier than the candidates, as it consistently was there; medians are reported alongside in that situation.

5. Some repartitioned meshes crash the model

Eight of the partitions generated during this campaign cause the model to diverge: on NG5, MINCONN-family candidates at both 2,048 and 64 ranks and a slack-only candidate at 64; on DARS, MINCONN-family candidates at 2,048 ranks; and on DARS at 2,048 ranks one partition produced by the *unmodified shipped recipe* with a different random seed, which diverged within fourteen steps. The failure is always local: one region runs away while the global diagnostics remain normal for tens of steps.

Three properties of these failures matter for practice.

First, they carry no offline signature. In the NG5 group, the partitions that later diverged scored better than the survivors on every precomputed quantity we examined, including load imbalance, edge cut and the number of isolated nodes.

Second, short tests do not find them. A five-step run passes on a partition that diverges at step 71,




Figure 3: The three GPU settings re-timed on GH200 (hatched) against their A100 values (solid), with the GH200 node count under each pair. The GH200 values for fArc and DARS are medians, because the shipped partition was the noisiest configuration in those jobs and the best-of- N estimator is then biased; the fArc value is indistinguishable from zero at the measured noise level.

and the 150-step timing runs passed on a partition that diverged before step 3,000. The 3,000-step requirement of Section 3 caught every failure listed above; nothing shorter would have.

Third, the failure belongs to the individual partition, not to the settings that produced it. On both meshes where a winning candidate diverged, regenerating it with the same settings and a different METIS seed gave a partition that passed — the full 3,000-step test on DARS, the 150-step timing run on NG5 — in one case with a worse edge cut than the partition that had failed. Conversely, the settings of the DARS GPU winner produce a sound partition at 64 ranks and a fatally unstable one at 2,048.

Since the shipped recipe itself produced one of the eight, the shipped partitions in production differ from the failures only in having been exercised. The practical conclusion is that the stability test is a property of adopting *any* new partition, not a surcharge on the new options.

6. Accuracy: the size of the disturbance, and what it means

Every change of decomposition perturbs the solution; the useful questions are by how much, on what scale, and through what mechanism. The scale bar is the spread of seed-only re-rolls of the shipped recipe (Section 3) — a disturbance no larger than that is one nobody would ever audit, since it is produced by regenerating today’s partition with a different random seed. The controls themselves span a factor 1.05 on DARS and 3.9 on CORE2 at 864 ranks, so no fixed rms threshold could replace the per-mesh spread. Figure 4 shows every surviving candidate’s disturbance per field on that scale; four grades occur in the campaign.

Inside or below the spread. Most recommended candidates: at CORE2, 4 GPU ranks, the winner lies below every control on all three fields; KaMinPar at DARS 2,048 lies below every control on temperature and salinity; Mt-KaHyPar at fArc and both CORE2 512 settings sit inside. Nothing to decide.

An SSH-only excursion with a mechanism. The DARS 64-GPU winner exceeds the spread in the SSH rms alone (+11%). The barotropic solver stops at a relative residual, so a partition on which it converges faster — which the MINCONN family produces — stops earlier and lands measurably elsewhere *within the same requested tolerance*: at DARS all three candidates converge faster than all three controls, and at CORE2 the candidate with the fewest iterations has the largest SSH difference. Both runs satisfy the tolerance asked of the solver; the difference lives inside the solver’s own tolerance ball, and tightening the tolerance would shrink it. This is a numerics artefact, not a physics bias.

An SSH-only excursion without one. The NG5 64-GPU winner shows the same signature (+8%),




Figure 4: Twenty-step rms difference from the reference for each candidate, by field, relative to the largest value among that point’s seed-control repartitionings (dashed line; the shaded region is inside the spread). Filled markers exceed the spread. Every candidate shown passed the 3,000-step stability screen. The DARS 2,048-rank envelope uses the four surviving controls (the fifth is the step-14 blowup of Section 5); the Hilbert-renumbered CORE2 candidate is compared under its own numbering and is omitted. All values are read directly from the accuracy-run logs; job ids are recorded in the campaign log.

but its solver iteration trace matches the controls exactly, ruling that mechanism out. Small and confined to one diagnostic field, but it stays labelled unexplained.

Temperature or salinity excursions. The -19.7% DARS arm has temperature rms $+13\%$ above the spread and a temperature *maximum* 1.8–2.5 times the control class; the fArc GPU candidates exceed a four-control temperature envelope by 18–33%; KaMinPar at CORE2 864 has salinity rms $+24\%$ (with temperature below every control). These are the disturbances that warrant attention: temperature and salinity carry the model’s long-term state, and a local maximum excursion is the same signature family as the local runaways that destroy partitions outright (Section 5). For such a candidate the honest path to production is a long twin integration against the shipped partition, not a threshold argument.

All numbers in this section are twenty-step cold-start divergences: they measure whether the partition change is *detectable* against ordinary repartitioning noise, not its climate impact. The inexpensive next tier of evidence is an end-state comparison after 3,000 steps against three controls; the definitive one is a multi-year twin run.

7. Recommendations

For FESOM2 upstream, in order of value:

1. Call `METIS_PartGraphKway` in `fort_part.c`, so that MINCONN, CONTIG and the volume objective reach METIS, and expose the partitioner options as run-time switches rather than compile-time constants.
2. Adopt the stability requirement together with the options: every newly generated partition, whatever produced it, runs 3,000 steps at its target rank count and the mesh’s cold-start time step before production use, and a failure is answered by regenerating with a new seed.
3. Rebuild the isolated-node repair in `check_partitioning`, which currently considers a single neighbour as reassignment candidate, and update METIS to 5.2.1.

On the machines measured here, the settings to reach for are MINCONN for GPU runs, with CONTIG and slack decided per configuration by a timing run, and an external engine (KaMinPar or Mt-KaHyPar, $w = 100 + n_{lev}$) or UFACTOR=30 for CPU runs. Hilbert renumbering is worthwhile on CORE2-class meshes where re-establishing the references has been accepted. Offline partition quality metrics are useful for designing candidates and useless for accepting them; acceptance is the stability test plus the disturbance report of Section 6, with a long twin run for any candidate whose disturbance reaches the temperature–salinity grade.

8. Status

Every measurement has run. The accuracy criterion was revised at campaign close: instead of a certified/not-certified verdict, each recommendation carries its measured disturbance and an interpretation (Section 6). Under that reporting the DARS and NG5 64-GPU winners (−14.3 and −9.8%, SSH-only disturbances) are ordinary recommendations, while the −19.7% DARS arm and the fArc 16-GPU point carry temperature-grade disturbances and want a long twin run before production. Four winners are already packaged as evidence-checked mesh directories under `mesh_m11_certified/`; the upstream pull request is the remaining engineering step. Whether an interconnect with working GPU–direct communication prices messages like the A100 fabric or like the GH200 testbed is open, and JUPITER should answer it by measurement.

Evidence trail: docs/PARTITIONING_M11.md (Findings 1–45, with job ids for every number quoted here); operational summary: docs/M11_RECOMMENDATION.md; full table of timing runs: scripts/m11_harvest_races.py -best.